

Day 18: The Chain Rule: The Computational Engine of Deep Learning

How a 300-Year-Old Mathematical Rule Powers Modern Artificial Intelligence

1. The Fundamental Principle: Composing Derivatives

The chain rule is a fundamental theorem in calculus that allows us to compute the derivative of a composite function. If $y = f(g(x))$, then:

$$\frac{dy}{dx} = \frac{dy}{dg} \cdot \frac{dg}{dx}$$

For deeper compositions $y = f(g(h(x)))$, this extends to:

$$\frac{dy}{dx} = \frac{dy}{df} \cdot \frac{df}{dg} \cdot \frac{dg}{dh} \cdot \frac{dh}{dx}$$

This simple multiplicative principle becomes extraordinarily powerful when applied to the deeply nested functions that represent neural networks.

Rohit Sanwariya

2. A Detailed Example: From Simple to Complex

Let's trace the chain rule through a more realistic example that mirrors a neural network layer.

Step-by-Step Chain Rule Application

Suppose we have:

$$z = h(x) = 2x + 1, \quad a = g(z) = z^2, \quad y = f(a) = \sin(a)$$

So the full composition is: $y = \sin((2x + 1)^2)$

To find $\frac{dy}{dx}$:

$$\begin{aligned} \frac{dy}{dx} &= \frac{dy}{da} \cdot \frac{da}{dz} \cdot \frac{dz}{dx} \\ &= \cos(a) \cdot 2z \cdot 2 \\ &= \cos((2x + 1)^2) \cdot 2(2x + 1) \cdot 2 \\ &= 4(2x + 1) \cdot \cos((2x + 1)^2) \end{aligned}$$

Notice how each function only needs to know its **local** derivative, which then gets multiplied with derivatives from other layers.

3. Computational Graphs: Visualizing the Chain Rule

We can represent complex functions as computational graphs where nodes are operations and edges show data flow:



The chain rule tells us that gradients flow backwards through this graph, with each node multiplying its local gradient by incoming gradients from downstream.

4. Backpropagation: The Chain Rule at Scale

Backpropagation is simply the systematic application of the chain rule to compute gradients in neural networks. The algorithm works in two phases:

The Backpropagation Algorithm

1. **Forward Pass:** Compute the network's output from input to output, storing all intermediate values.
2. **Backward Pass:**
 - Compute the gradient of the loss with respect to the output
 - Propagate this gradient backward through the network
 - At each layer, multiply by the local derivative of that layer's operation
 - Accumulate parameter gradients as you go

This approach is efficient because it avoids recomputing shared subexpressions—each node's gradient is computed exactly once.

5. A Concrete Neural Network Example

Consider a simple network with one hidden layer:

$$z = W_1x + b_1, \quad a = \sigma(z), \quad \hat{y} = W_2a + b_2, \quad L = \frac{1}{2}(\hat{y} - y)^2$$

Using the chain rule, we compute the gradient for W_1 :

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial W_1} = (\hat{y} - y) \cdot W_2 \cdot \sigma'(z) \cdot x$$

Each term is a local derivative that depends only on the operation at that layer.

6. Why This Matters: The Practical Implications

How the Chain Rule Enables Modern AI

- **Efficiency:** Backpropagation computes all gradients in one forward and one backward pass ($O(n)$ time) instead of $O(n^2)$ for naive methods.
- **Modularity:** Each layer only needs to implement its forward pass and local gradient computation. This allows for flexible architecture design.
- **Automatic Differentiation:** Frameworks like TensorFlow and PyTorch automatically construct the computational graph and apply the chain rule.
- **Memory Tradeoff:** Backpropagation requires storing intermediate values from the forward pass (the "activation memory" problem in large models).

7. Beyond Basic Backpropagation: Modern Extensions

The basic chain rule concept has been extended to handle more complex scenarios:

- **Implicit Layers:** Layers where the output is defined implicitly (e.g., through an equation) require specialized gradient computation.
- **Reversible Networks:** Architectures that don't need to store intermediate activations by recomputing them during the backward pass.
- **Higher-Order Derivatives:** Second-order optimization methods need derivatives of derivatives, leading to complex chain rule applications.
- **Meta-Learning:** Calculating gradients through gradient descent steps requires careful application of the chain rule.

8. Historical Context and Significance

The chain rule was formalized by Gottfried Wilhelm Leibniz in the 17th century, but its application to neural networks only gained prominence in the 1980s and 1990s. The 1986 paper "Learning representations by back-propagating errors" by Rumelhart, Hinton, and Williams was pivotal in popularizing backpropagation.

What's remarkable is that a mathematical principle discovered centuries before the computer revolution would become the computational foundation for one of the most transformative technologies of the 21st century.

Key Takeaway

The chain rule is not just a mathematical curiosity—it is the fundamental algorithm that makes training deep neural networks computationally feasible. By allowing us to break down the complex gradient computation of nested functions into a series of simple local operations, it enables the efficient training of models with millions or even billions of parameters. Understanding the chain rule is essential for anyone who wants to truly understand how deep learning works, debug training issues, or develop new architectural innovations.

Rohit Sanwariya