

Day 9: Poisson Distribution — Modeling Rare Events and Count Data

Understanding Arrival Processes and Discrete Event Modeling in AI

1. The Fundamental Question

How do we mathematically model events that occur randomly over time or space, particularly when they're rare but important, and how does this enable AI systems to understand arrival processes, count data, and temporal patterns?

The Poisson distribution provides the mathematical framework for modeling count data and rare events, forming the statistical foundation for time series analysis, natural language processing, and event prediction in artificial intelligence systems.

2. Mathematical Foundations

2.1. Poisson Distribution Definition

A discrete random variable X follows a Poisson distribution with rate parameter $\lambda > 0$ if:

$$P(X = k) = \frac{\lambda^k e^{-\lambda}}{k!}, \quad k = 0, 1, 2, \dots$$

2.2. Key Properties and Moments

$$\text{Mean: } E[X] = \lambda$$

$$\text{Variance: } \text{Var}(X) = \lambda$$

$$\text{Skewness: } \gamma_1 = \frac{1}{\sqrt{\lambda}}$$

$$\text{Kurtosis: } \gamma_2 = \frac{1}{\lambda}$$

$$\text{MGF: } M_X(t) = e^{\lambda(e^t - 1)}$$

2.3. Poisson Process Foundations

Poisson Process Properties

- **Stationary Increments:** Probability depends only on interval length, not starting point
- **Independent Increments:** Non-overlapping intervals are independent
- **Orderliness:** Multiple events cannot occur simultaneously
- **Exponential Inter-arrival:** Time between events follows exponential distribution
- **Additive Property:** Sum of independent Poisson variables is Poisson

3. Comprehensive Examples and Analysis

3.1. Basic Poisson Example: Customer Arrivals

Service System Modeling

- Call center receives average of 5 calls per hour ($\lambda = 5$)
- Model number of calls in one hour: $X \sim \text{Poisson}(5)$
- Applications: Queueing systems, service capacity planning, resource allocation

Probability Calculations:

$$P(X = 0) = \frac{5^0 e^{-5}}{0!} = e^{-5} \approx 0.0067$$

$$P(X = 1) = \frac{5^1 e^{-5}}{1!} = 5e^{-5} \approx 0.0337$$

$$P(X = 2) = \frac{5^2 e^{-5}}{2!} = 12.5e^{-5} \approx 0.0842$$

$$P(X \leq 2) = P(X = 0) + P(X = 1) + P(X = 2) \approx 0.1246$$

$$P(X \geq 3) = 1 - P(X \leq 2) \approx 0.8754$$

Interpretation: There's an 87.5% chance of receiving 3 or more calls per hour.

3.2. Advanced Example: Web Server Requests

High-Traffic System Analysis

- Web server handles average of 1000 requests per minute ($\lambda = 1000$)
- Use normal approximation for large λ
- Calculate probability of exceeding capacity (1200 requests/minute)

Normal Approximation:

$$\begin{aligned}X &\sim \text{Poisson}(1000) \approx \mathcal{N}(1000, 1000) \\P(X > 1200) &= 1 - P\left(Z \leq \frac{1200 - 1000}{\sqrt{1000}}\right) \\&= 1 - P(Z \leq 6.324) \approx 1 - 1 = 0\end{aligned}$$

With Continuity Correction:

$$\begin{aligned}P(X > 1200) &= P(X \geq 1200.5) \\&= 1 - P\left(Z \leq \frac{1200.5 - 1000}{\sqrt{1000}}\right) \\&= 1 - P(Z \leq 6.344) \approx 0\end{aligned}$$

Interpretation: Probability of exceeding 1200 requests is negligible under normal operation.

3.3. Real-World Application: Network Anomaly Detection

Cybersecurity Application

- Normal network: 10 failed login attempts per hour ($\lambda = 10$)
- Detect brute force attacks: threshold at 25 attempts/hour
- Calculate false positive rate and detection probability

Statistical Analysis:

$$\begin{aligned}\text{False Positive Rate} &= P(X \geq 25 | \lambda = 10) \\&= 1 - \sum_{k=0}^{24} \frac{10^k e^{-10}}{k!} \approx 1.13 \times 10^{-5}\end{aligned}$$

If attack ($\lambda = 50$): $P(X \geq 25) \approx 0.99998$

Performance: System achieves 99.998% detection rate with very low false positives.

4. Why Poisson Distribution is Fundamental to AI

Critical Applications in Machine Learning

- **Natural Language Processing:** Word counts, document topic modeling
- **Computer Vision:** Photon counting, particle detection
- **Time Series Analysis:** Event counting, arrival process modeling
- **Anomaly Detection:** Rare event modeling and outlier identification
- **Reinforcement Learning:** Event-based reward modeling
- **Network Analysis:** Graph degree distributions, connection patterns

5. Real-World AI Applications

5.1. Natural Language Processing

- **Word Frequency:** Poisson modeling of word occurrences in documents
- **Topic Modeling:** Poisson factor analysis for document generation
- **Information Retrieval:** Term frequency modeling in search engines
- **Performance:** Poisson-based models achieve 85-92% accuracy in text classification

5.2. Computer Vision and Sensor Systems

- **Low-Light Imaging:** Photon counting statistics in cameras
- **Particle Detection:** Radioactive decay counting in medical imaging
- **Event Cameras:** Asynchronous pixel-level event detection
- **Performance:** Poisson noise models improve image quality by 20-40% in low-light conditions

5.3. Network and System Monitoring

- **Network Traffic:** Modeling packet arrivals and network congestion
- **Server Monitoring:** Request arrival patterns and load balancing
- **Financial Transactions:** Trade arrival modeling in high-frequency trading
- **Performance:** Anomaly detection systems achieve 95%+ precision in production environments

6. Implementation in Modern AI Frameworks

6.1. PyTorch Implementation

```
import torch
import torch.distributions as dist
import matplotlib.pyplot as plt

# Basic Poisson distribution
lambda_val = 5.0
poisson = dist.Poisson(lambda_val)

# PMF calculation
k_values = torch.arange(0, 15)
pmf = torch.exp(poisson.log_prob(k_values))
print("PMF values:", pmf)

# Sampling and empirical verification
samples = poisson.sample((10000,))
empirical_mean = samples.float().mean()
empirical_var = samples.float().var()

print(f"Theoretical: E[X]={lambda_val}, Var(X)={lambda_val}")
print(f"Empirical: E[X]={empirical_mean:.3f}, Var(X)={empirical_var:.3f}")

# Poisson regression for count data
class PoissonRegression(torch.nn.Module):
    def __init__(self, input_dim):
        super().__init__()
        self.linear = torch.nn.Linear(input_dim, 1)
```

```

def forward(self, x):
    # Ensure positive rate parameter
    rate = torch.exp(self.linear(x))
    return dist.Poisson(rate)

# Application: Text classification with bag-of-words
def poisson_naive_bayes(train_features, train_labels):
    """Poisson Naive Bayes for text classification"""
    n_classes = train_labels.max() + 1
    n_features = train_features.shape[1]

    # Calculate class priors
    class_priors = torch.zeros(n_classes)
    class_lambdas = torch.zeros(n_classes, n_features)

    for c in range(n_classes):
        class_mask = (train_labels == c)
        class_priors[c] = class_mask.float().mean()
        class_lambdas[c] = train_features[class_mask].mean(dim=0)

    return class_priors, class_lambdas

```

6.2. Advanced Poisson Modeling

```

import numpy as np
from scipy import stats

def compound_poisson_process(lambda_val, claim_distribution, T):
    """Simulate compound Poisson process (e.g., insurance claims)"""
    # Number of events in [0, T]
    n_events = stats.poisson(lambda_val * T).rvs()

    # Event times uniformly distributed in [0, T]
    event_times = np.sort(T * np.random.uniform(0, 1, n_events))

    # Claim sizes
    claim_sizes = claim_distribution.rvs(n_events)

    return event_times, claim_sizes

def poisson_regression_glm(X, y):

```

```

"""Generalized Linear Model with Poisson response"""
from sklearn.linear_model import PoissonRegressor

model = PoissonRegressor(alpha=0.1, max_iter=1000)
model.fit(X, y)

# Predict expected counts
y_pred = model.predict(X)

# Calculate deviance residuals
residuals = np.sign(y - y_pred) * np.sqrt(
    2 * (y * np.log(y/y_pred) - (y - y_pred))
)

return model, y_pred, residuals

# Zero-inflated Poisson for excess zeros
def zero_inflated_poisson_log_prob(y, lambda_val, pi):
    """Log probability for zero-inflated Poisson"""
    # pi: probability of structural zero
    log_prob = torch.zeros_like(y, dtype=torch.float)

    # For zero observations
    zero_mask = (y == 0)
    log_prob[zero_mask] = torch.log(
        pi[zero_mask] + (1 - pi[zero_mask]) * torch.exp(-lambda_val[zero_mask])
    )

    # For positive observations
    pos_mask = (y > 0)
    log_prob[pos_mask] = torch.log(1 - pi[pos_mask]) + \
        dist.Poisson(lambda_val[pos_mask]).log_prob(y[pos_mask])

    return log_prob

```

7. Advanced Concepts and Extensions

7.1. Related Distributions

- **Exponential Distribution:** Inter-arrival times in Poisson process

- **Gamma Distribution:** Waiting time for k-th event
- **Negative Binomial:** Overdispersed count data (variance $>$ mean)
- **Compound Poisson:** Sum of random variables with Poisson number of terms
- **Zero-Inflated Poisson:** Excess zeros in count data

7.2. Poisson Process Extensions

- **Homogeneous Poisson Process:** Constant rate λ
- **Non-homogeneous Poisson Process:** Time-varying rate $\lambda(t)$
- **Spatial Poisson Process:** Events distributed in space rather than time
- **Cox Process:** Doubly stochastic Poisson process with random intensity
- **Hawkes Process:** Self-exciting process where events increase future intensity

7.3. Statistical Inference and Testing

- **Maximum Likelihood:** $\hat{\lambda} = \frac{1}{n} \sum_{i=1}^n x_i$
- **Goodness-of-Fit:** Chi-square test, dispersion test
- **Overdispersion:** Variance $>$ mean indicates negative binomial may be better
- **Zero Inflation:** Excess zeros indicate ZIP model may be appropriate

8. Performance Analysis and Empirical Validation

Poisson Modeling Best Practices

- **Mean-Variance Equality:** Check if sample variance \approx sample mean
- **Dispersion Test:** Use Fisher's dispersion test for overdispersion
- **Zero Inflation:** Compare observed vs expected zeros
- **Time Homogeneity:** Check if rate is constant over time
- **Independence:** Verify events are independent (autocorrelation tests)
- **Model Selection:** Use AIC/BIC for Poisson vs negative binomial

9. Practical Exercises for Mastery

Hands-On Poisson Analysis

1. **Basic Properties:** Derive mean, variance, and MGF for $\text{Poisson}(\lambda)$
2. **Additive Property:** Prove sum of independent Poissons is Poisson
3. **Normal Approximation:** Compare exact Poisson probabilities with normal approximation
4. **Process Simulation:** Implement homogeneous and non-homogeneous Poisson processes
5. **Real Data Analysis:** Fit Poisson distribution to real count data (web traffic, word counts)
6. **Goodness-of-Fit:** Implement chi-square test for Poisson distribution
7. **Regression Modeling:** Build Poisson regression model for count data prediction

10. Common Pitfalls and Best Practices

Poisson Modeling Guidelines

- **Overdispersion:** Use negative binomial if variance \neq mean
- **Zero Inflation:** Consider zero-inflated models for excess zeros
- **Rate Variation:** Check if rate is constant over time/space
- **Independence:** Ensure events are independent (no clustering)
- **Large λ :** Use normal approximation for $\lambda > 20$
- **Small Counts:** Use exact methods rather than approximations
- **Model Diagnostics:** Always check residuals and goodness-of-fit

11. Historical Context and Modern Impact

11.1. Historical Development

- **1837:** Siméon Denis Poisson introduces distribution for jury decisions

- **Late 19th:** Ladislaus Bortkiewicz applies to Prussian cavalry deaths
- **Early 20th:** Applications in radioactivity (Rutherford Geiger)
- **Mid 20th:** Queueing theory and telecommunications applications
- **Late 20th:** Generalized linear models and regression applications

11.2. Modern AI Applications

- **Deep Learning:** Poisson noise models in variational autoencoders
- **Natural Language Processing:** Poisson factorization for topic modeling
- **Computer Vision:** Photon-limited imaging and denoising
- **Reinforcement Learning:** Poisson processes for event-based rewards
- **Time Series:** Point process models for irregularly spaced events

12. Key Insight: The Universal Model for Rare Events

The Poisson distribution provides a powerful mathematical framework for understanding and modeling count data and rare events in AI systems:

- **Memoryless Property:** Future depends only on present, not past
- **Scale Invariance:** Works across different time/space scales
- **Computational Efficiency:** Simple mathematical form enables fast computation
- **Theoretical Foundation:** Well-understood properties support rigorous analysis
- **Practical Versatility:** Applicable across diverse domains and applications

The elegance of the Poisson distribution lies in its ability to:

- **Model Rare Events:** Capture low-probability but high-impact occurrences
- **Handle Count Data:** Naturally represent discrete event counts
- **Support Temporal Analysis:** Model arrival processes and time-based patterns
- **Enable Spatial Modeling:** Extend to spatial point processes
- **Form Building Blocks:** Serve as foundation for more complex distributions

Mastering the Poisson distribution is essential for building AI systems that can effectively model count data, understand temporal patterns, detect anomalies in event streams, and make predictions about rare but important events—skills that are fundamental to modern AI applications in cybersecurity, natural language processing, computer vision, and system monitoring.

Next: Exponential and Gamma Distributions — Modeling Waiting Times and Durations

Tomorrow we'll explore the Exponential and Gamma distributions and their applications in modeling waiting times, durations, and time-to-event data—essential for survival analysis, reliability engineering, and temporal modeling in AI systems.

Rohit Sanwariya